

SAURABH BHAUSAHEB ZINJAD

480-913-5544 [linkedin.com/in/saurabhzinjad](https://www.linkedin.com/in/saurabhzinjad) saurabhzinjad@gmail.com github.com/Ztrimus

Versatile engineer with 6 years of experience developing production-ready ML solutions and full stack applications. Managing end-to-end pipelines, from design to deployment and optimization in cloud environments. Passionate about advancing AI infrastructure, fine-tuning foundation models, and delivering scalable, cost-effective solutions globally.

SKILLS

Languages: Python, JavaScript/TypeScript, C#, C++, SQL, Java, Bash, HTML, CSS
ML/AI: PyTorch, TensorFlow, scikit-learn, LangChain, LangGraph, OpenCV, Pandas, MLflow, Databricks
GenAI & Systems: RAG, LLM evaluation, fine-tuning, prompt engineering, agentic AI design, MLOps, observability, Kafka
Cloud/DevOps: AWS (Bedrock, Lambda, S3, CloudWatch API Gateway, Cognito), GCP, Azure, Docker, Kubernetes, Git
Full-Stack: React, Node.js, Flask, FastAPI, .NET, PostgreSQL, MySQL, MongoDB, Terraform

WORK EXPERIENCE

- AI/ML Engineer II** — *University of Phoenix, AZ, USA* Feb 2025 – Present
- Built GenAI + explainability tools (LangGraph + OpenAI) to process unstructured data and auto-explain complex academic/financial calcs, enhance advisor support capacity by ~ 6x
 - Integrated organizational knowledge base with LLMs, building pipelines to extract code logic into step-by-step explanations, accelerating onboarding and improving developer productivity
- AI Research Intern** — *University of Phoenix, AZ, USA* May 2024 – Aug 2024
- LLM summary dashboard with chain-of-thought reasoning integrated to Salesforce via Kafka; advisor efficiency ↑ 40%
 - Text-to-Visual API service converting text to interactive diagrams; integrated with Slack/Teams/CRMs; documentation speed ↑
- Senior Machine Learning Engineer** — *Tiger Analytics, Bangalore, India* June 2022 – July 2023
- Trained & tuned predictive ML models and scalable Data pipelines (Python, PySpark, Databricks) achieving 18% client revenue
 - Architected a multi-cloud model monitoring platform for lifecycle tracking/alerts; supported 10+ client adoptions
- Software Engineer** — *Winjit Technologies, Nashik, India* January 2020 – June 2022
- Engineered 10+ RESTful APIs for real-time AI data processing; built 30+ low latency responsive UI React features
 - Optimized large-scale MySQL, PostgreSQL databases, accelerating data access for AI model training and real-time inference
 - Innovated data-driven UI framework with dynamic form reconfiguration, customizable CSS, reduced development time by 8x
- Deep Learning Engineer** — *Automation Teknix, Pune, India* September 2019 – January 2020
- Developed lightweight object recognition engine with SSD & MobileNetV2, reducing computational cost by 30% for mobile use
 - Streamlined CNN model using TensorFlow, OpenCV to improve image processing, reducing survey error in production by 22%

EDUCATION

- Arizona State University, Tempe, USA** August 2023 – May 2025
Masters of Science in Computer Science - **Thesis**
- Pune Institute of Computer Technology, Pune, India** July 2015 – June 2019
Bachelor of Engineering in Electronics and Telecommunications

PUBLICATION *first authored*

- Can Typos Cause Harm? The Impact of Imperfect Input on LLM Safety** Sep 2024 – Aug 2025
- Systematic study of LLM safety under prompt perturbations; introduced safety flip rate and evaluated 5 models across 14 perturbation types & 11 harm categories—minor typos/paraphrases flip safety in 30% of cases
 - Released open-source framework & augmentation pipeline for perturbation-aware safety evaluation and robust alignment tools
- LLM-facilitated Pipeline for Personalized Resume Generation** Aug 2023 – July 2024
- Authored & published research paper at ACM's SIGIR conference on ML pipeline optimized for user personalization using LLM
 - Formulated novel user personalization metrics, setting performance benchmarks for AI-enhanced resume evaluation
 - Released open-source Python library integrating information retrieval leveraging web scraping, used by 1000+ users

PROJECTS

- Multimodal AI System for Media Search — AWS + CLIP + vector DB** October 2024 – October 2024
- Won Big Data Prize for content-based search across video/audio/images; stored in Postgres delivered 92% faster media
- Forest fire Detection using IoT Sensor Data** September 2021 – January 2022
- Implemented a TabNet classifier model with 98.7% accuracy on edge devices using AWS, TinyML, Docker, Redis, and Celery
 - Utilized SMOTE and hyperparameter tuning to balance data and optimize recall, reducing errors detecting fire in 7 mins